

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
credit_df = pd.read_csv('credit_card.csv')
```

```
credit_df.head()
```

	Time	V1	V2	V3	V4	V5	V6
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921

	V8	V9	...	V21	V22	V23	V24
0	0.098698	0.363787	...	-0.018307	0.277838	-0.110474	0.066928
1	0.085102	-0.255425	...	-0.225775	-0.638672	0.101288	-0.339846
2	0.247676	-1.514654	...	0.247998	0.771679	0.909412	-0.689281
3	0.377436	-1.387024	...	-0.108300	0.005274	-0.190321	-1.175575
4	-0.270533	0.817739	...	-0.009431	0.798278	-0.137458	0.141267

	V26	V27	V28	Amount	Class
0	-0.189115	0.133558	-0.021053	149.62	0
1	0.125895	-0.008983	0.014724	2.69	0
2	-0.139097	-0.055353	-0.059752	378.66	0
3	-0.221929	0.062723	0.061458	123.50	0
4	0.502292	0.219422	0.215153	69.99	0

```
[5 rows x 31 columns]
```

#The credit card data is very sensitive and confidential and thus we have numerical values are given

```
credit_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 284807 entries, 0 to 284806
```

```
Data columns (total 31 columns):
#      Column      Non-Null Count  Dtype
---  -
0      Time        284807 non-null  float64
1      V1           284807 non-null  float64
2      V2           284807 non-null  float64
3      V3           284807 non-null  float64
4      V4           284807 non-null  float64
5      V5           284807 non-null  float64
6      V6           284807 non-null  float64
7      V7           284807 non-null  float64
8      V8           284807 non-null  float64
9      V9           284807 non-null  float64
10     V10          284807 non-null  float64
11     V11          284807 non-null  float64
12     V12          284807 non-null  float64
13     V13          284807 non-null  float64
14     V14          284807 non-null  float64
15     V15          284807 non-null  float64
16     V16          284807 non-null  float64
17     V17          284807 non-null  float64
18     V18          284807 non-null  float64
19     V19          284807 non-null  float64
20     V20          284807 non-null  float64
21     V21          284807 non-null  float64
22     V22          284807 non-null  float64
23     V23          284807 non-null  float64
24     V24          284807 non-null  float64
25     V25          284807 non-null  float64
26     V26          284807 non-null  float64
27     V27          284807 non-null  float64
28     V28          284807 non-null  float64
29     Amount       284807 non-null  float64
30     Class        284807 non-null  int64
```

```
dtypes: float64(30), int64(1)
```

```
memory usage: 67.4 MB
```

#Most of the values are in float. So no string/object data is present.

```
credit_df.isnull().sum()
```

```
Time      0
V1        0
V2        0
V3        0
V4        0
V5        0
V6        0
V7        0
V8        0
```

```

V9      0
V10     0
V11     0
V12     0
V13     0
V14     0
V15     0
V16     0
V17     0
V18     0
V19     0
V20     0
V21     0
V22     0
V23     0
V24     0
V25     0
V26     0
V27     0
V28     0
Amount  0
Class   0
dtype: int64

```

#Therefore there are no null values in the given data

```
credit_df.describe()
```

	Time	V1	V2	V3
V4 \				
count	284807.000000	2.848070e+05	2.848070e+05	2.848070e+05
mean	94813.859575	1.759061e-12	-8.251130e-13	-9.654937e-13
std	47488.145955	1.958696e+00	1.651309e+00	1.516255e+00
min	0.000000	-5.640751e+01	-7.271573e+01	-4.832559e+01
25%	54201.500000	-9.203734e-01	-5.985499e-01	-8.903648e-01
50%	84692.000000	1.810880e-02	6.548556e-02	1.798463e-01
75%	139320.500000	1.315642e+00	8.037239e-01	1.027196e+00
max	172792.000000	2.454930e+00	2.205773e+01	9.382558e+00
	V5	V6	V7	V8
V9 \				
count	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05

```

2.848070e+05
mean    1.649999e-13  4.248366e-13 -3.054600e-13  8.777971e-14 -
1.179749e-12
std      1.380247e+00  1.332271e+00  1.237094e+00  1.194353e+00
1.098632e+00
min     -1.137433e+02 -2.616051e+01 -4.355724e+01 -7.321672e+01 -
1.343407e+01
25%     -6.915971e-01 -7.682956e-01 -5.540759e-01 -2.086297e-01 -
6.430976e-01
50%     -5.433583e-02 -2.741871e-01  4.010308e-02  2.235804e-02 -
5.142873e-02
75%      6.119264e-01  3.985649e-01  5.704361e-01  3.273459e-01
5.971390e-01
max      3.480167e+01  7.330163e+01  1.205895e+02  2.000721e+01
1.559499e+01

```

	...	V21	V22	V23	V24	\
count	...	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	
mean	...	-3.405756e-13	-5.723197e-13	-9.725856e-13	1.464150e-12	
std	...	7.345240e-01	7.257016e-01	6.244603e-01	6.056471e-01	
min	...	-3.483038e+01	-1.093314e+01	-4.480774e+01	-2.836627e+00	
25%	...	-2.283949e-01	-5.423504e-01	-1.618463e-01	-3.545861e-01	
50%	...	-2.945017e-02	6.781943e-03	-1.119293e-02	4.097606e-02	
75%	...	1.863772e-01	5.285536e-01	1.476421e-01	4.395266e-01	
max	...	2.720284e+01	1.050309e+01	2.252841e+01	4.584549e+00	

	V25	V26	V27	V28
Amount \				
count	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05
284807.000000				
mean	-6.987102e-13	-5.617874e-13	3.332082e-12	-3.518874e-12
88.349619				
std	5.212781e-01	4.822270e-01	4.036325e-01	3.300833e-01
250.120109				
min	-1.029540e+01	-2.604551e+00	-2.256568e+01	-1.543008e+01
0.000000				
25%	-3.171451e-01	-3.269839e-01	-7.083953e-02	-5.295979e-02
5.600000				
50%	1.659350e-02	-5.213911e-02	1.342146e-03	1.124383e-02
22.000000				
75%	3.507156e-01	2.409522e-01	9.104512e-02	7.827995e-02
77.165000				
max	7.519589e+00	3.517346e+00	3.161220e+01	3.384781e+01
25691.160000				

	Class
count	284807.000000
mean	0.001727
std	0.041527
min	0.000000

```

25%      0.000000
50%      0.000000
75%      0.000000
max       1.000000

```

```
[8 rows x 31 columns]
```

#There are many outliers which we will try to reduce using the z_score method

```
credit_df['Class'].value_counts()
```

```

Class
0      284315
1         492
Name: count, dtype: int64

```

Here Label 0 represents normal transactions and Label 1 represents fraudulent transactions.

Also the number of normal transactions is far greater than the fraudulent ones because of which during the training and testing of the model, there can be development of biasness. This problem can be handled by using the Under Sampling further in this project

This is an unbalanced dataset

#Lets separate the two classes dataset

```

normal = credit_df[credit_df['Class'] == 0]
fraudulent = credit_df[credit_df['Class'] == 1]

```

```
normal.shape, fraudulent.shape
```

```
((284315, 31), (492, 31))
```

```
normal = normal.sample(n = 492)
```

```
new_df = pd.concat([normal, fraudulent])
```

```
new_df.head(), new_df.shape
```

```

(
      Time      V1      V2      V3      V4      V5
V6 \
86464    61247.0 -0.388777  0.794784  1.685546  0.916746  0.231280 -
0.288404
242179   151377.0  1.423943 -1.496709 -0.335678  0.834365 -1.071853
0.410509
67658    52642.0 -2.584287  0.396221  0.611749 -2.618310 -2.086440
0.749698
145724    87159.0  2.128941 -0.145630 -3.901839 -1.020558  3.172941
2.590078
15239    26595.0  1.239207 -0.501248  0.051282 -0.527667 -0.658764 -
0.273362

```

	V7	V8	V9	...	V21	V22
V23 \						
86464	0.803386	-0.252804	-0.778652	...	0.035846	0.204973 -
0.177384						
242179	-0.608726	0.192055	2.069181	...	-0.024910	-0.343123
0.101321						
67658	-1.639580	1.442094	-1.869689	...	-0.163833	-0.264298 -
0.264469						
145724	0.394099	0.410361	-0.031873	...	0.279982	0.900522 -
0.169174						
15239	-0.594065	0.076999	-0.899624	...	0.018870	0.004530 -
0.023754						

	V24	V25	V26	V27	V28	Amount
Class						
86464	0.327731	0.117845	-0.399662	-0.267336	-0.199676	38.97
0						
242179	0.659025	-0.418078	-0.532157	0.009470	0.020748	274.43
0						
67658	-1.009203	0.566115	-0.235755	-0.399279	-0.079796	60.00
0						
145724	0.799857	0.788816	0.480904	-0.084483	-0.096278	9.00
0						
15239	-0.040694	0.363785	-0.292766	0.028772	0.030351	44.99
0						

```
[5 rows x 31 columns],
(984, 31))
```

```
credit_df.groupby('Class').mean()
```

	Time	V1	V2	V3	V4	V5
\						
Class						
0	94838.202258	0.008258	-0.006271	0.012171	-0.007860	0.005453
1	80746.806911	-4.771948	3.623778	-7.033281	4.542029	-3.151225

	V6	V7	V8	V9	...	V20	V21
\							
Class					...		
0	0.002419	0.009637	-0.000987	0.004467	...	-0.000644	-0.001235
1	-1.397737	-5.568731	0.570636	-2.581123	...	0.372319	0.713588

	V22	V23	V24	V25	V26	V27
V28 \						
Class						
0	-0.000024	0.000070	0.000182	-0.000072	-0.000089	-0.000295
1	0.014049	-0.040308	-0.105130	0.041449	0.051648	0.170575

	Amount
Class	
0	88.291022
1	122.211321

[2 rows x 30 columns]

```
new_df.groupby('Class').mean()
```

	Time	V1	V2	V3	V4	V5
\						
Class						
0	92521.441057	0.036058	-0.116719	-0.020000	0.017799	-0.003164
1	80746.806911	-4.771948	3.623778	-7.033281	4.542029	-3.151225

	V6	V7	V8	V9	...	V20	V21
\							
Class					...		
0	0.070291	-0.068135	0.061256	-0.014633	...	0.020047	-0.024238
1	-1.397737	-5.568731	0.570636	-2.581123	...	0.372319	0.713588


	V22	V23	V24	V25	V26	V27
V28 \						
Class						
0	-0.031225	0.055564	0.013123	0.012087	-0.020505	-0.027695
1	0.014049	-0.040308	-0.105130	0.041449	0.051648	0.170575

	Amount
Class	
0	94.827297
1	122.211321

[2 rows x 30 columns]

#The nature of the dataset hasnt changed

```
plt.figure(figsize=(20, 10))
sns.heatmap(new_df.corr(), annot=True, cmap='coolwarm', fmt=".2f")
plt.show()
```



#Now lets split the dataset

```
corr_matrix = new_df.corr()
corr_matrix['Class'].sort_values(ascending = False)
```

```
Class      1.000000
V4         0.708342
V11        0.683403
V2         0.496818
V19        0.281222
V20        0.161317
V21        0.132174
V27        0.097089
V28        0.090502
V26        0.077539
Amount     0.054580
V8         0.052258
V25        0.022475
V22        0.019102
V23       -0.039721
V15       -0.053298
V13       -0.075723
```




```
V24      -0.103987
Time      -0.123582
V5        -0.372393
V6        -0.410638
V1        -0.432982
V18       -0.454542
V7        -0.470942
V9        -0.550945
V17       -0.560221
V3        -0.565386
V16       -0.600726
V10       -0.628843
V12       -0.681290
V14       -0.751085
Name: Class, dtype: float64
```

#V4,V11,V2 show a very strong positive correlation

#V14,V12,V10,V16,V3,V17,V9 show a very strong negative correlation

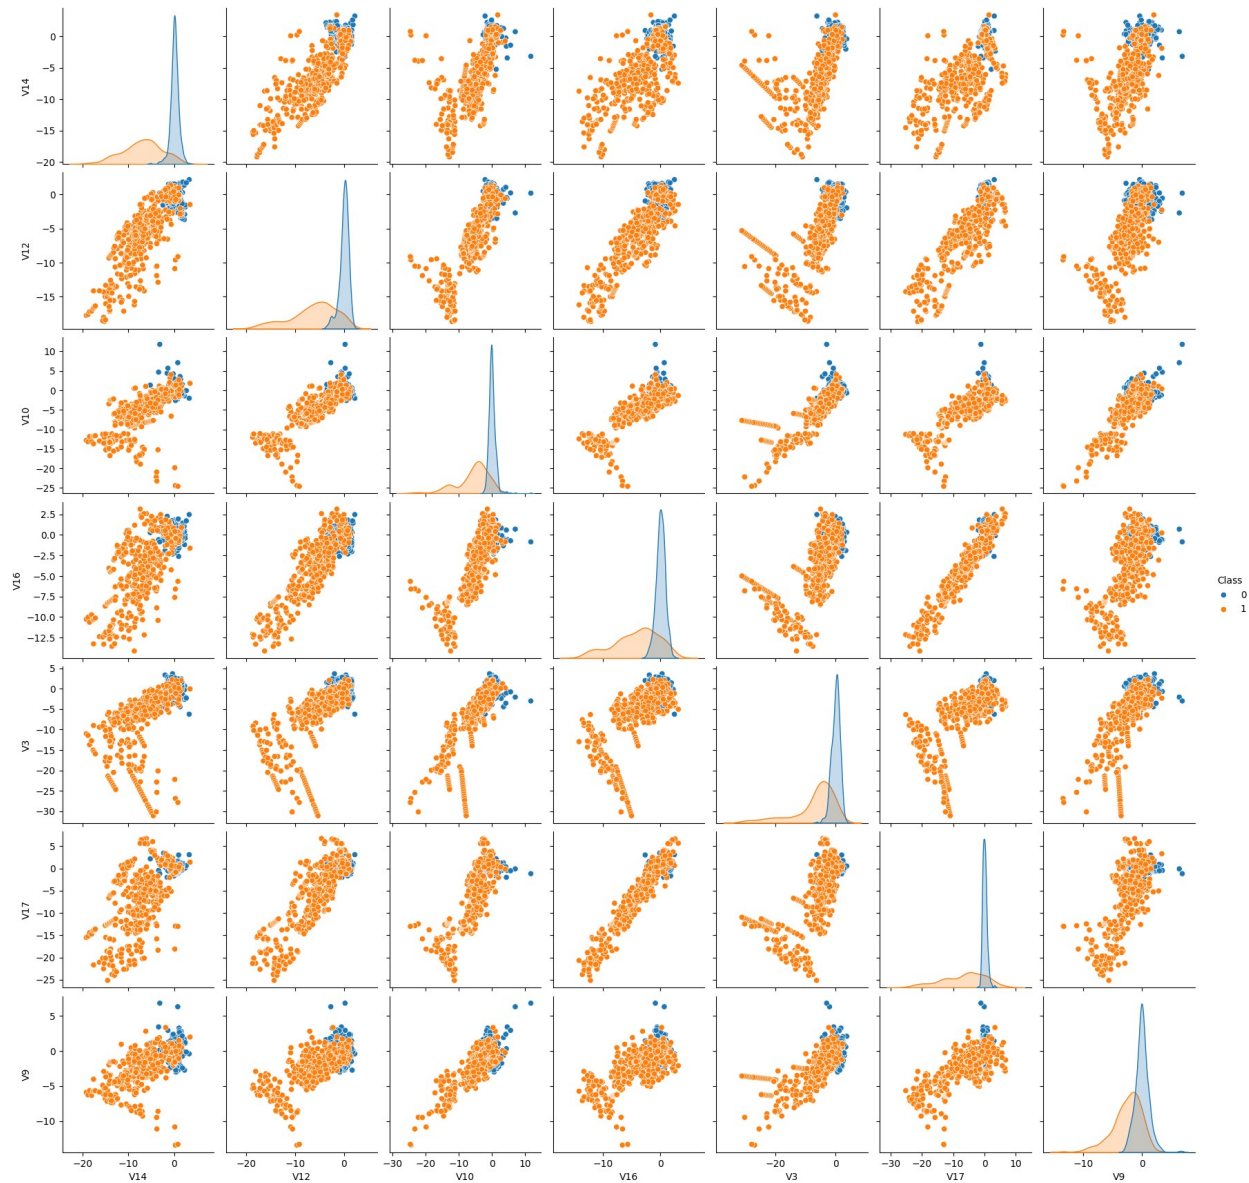
```
sns.pairplot(new_df[['V4', 'V11', 'V2', 'Class']], hue = 'Class')
plt.show()
```



#We know that we have the label as a categorical data. Thus the first model that comes to the mind is Logistic Regression.

```
sns.pairplot(new_df[['V14', 'V12', 'V10', 'V16', 'V3', 'V17', 'V9', 'Class']], hue = 'Class')
```

```
<seaborn.axisgrid.PairGrid at 0x23dd01c39b0>
```



#The variation in data seems to be fairly linear with respect to each other

#A model with the label and considering only the strongly related features can also be hold considerable

```
x = new_df.iloc[:, :-1]
y = new_df.Class
```

```
from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size =
0.2, random_state = 42)
```

```
from sklearn.linear_model import LogisticRegression
lr = LogisticRegression()
lr.fit(x_train,y_train)
```

C:\Users\gaikw\AppData\Local\Programs\Python\Python313\Lib\site-packages\sklearn\linear_model_logistic.py:465: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

```
LogisticRegression())
```

```
lr.score(x_test,y_test)
```

```
0.9238578680203046
```

```
lr.score(x_train,y_train)
```

```
0.9466327827191868
```

#Our model seems fairly considerable for usage but lets do some hyperparamter tuning to check if we can get some more accuracy

```
import warnings
```

```
warnings.filterwarnings('ignore')
```

```
df = {'penalty' : ['l1', 'l2', 'elasticnet', None], 'solver' :  
      ['lbfgs', 'liblinear', 'newton-cg', 'newton-cholesky', 'sag',  
      'saga'], 'multi_class' : ['auto', 'ovr', 'multinomial']}
```

```
from sklearn.model_selection import RandomizedSearchCV
```

```
rd = RandomizedSearchCV(LogisticRegression(),param_distributions = df,  
n_iter = 10)
```

```
rd.fit(x_train,y_train)
```

```
RandomizedSearchCV(estimator=LogisticRegression(),
```

```
                    param_distributions={'multi_class': ['auto', 'ovr',
```

```
'multinomial'],
```

```
                    'penalty': ['l1', 'l2',
```

```
'elasticnet',
```

```
                    None],
```

```
                    'solver': ['lbfgs',
```

```
'liblinear',
```

```
                    'newton-cg',
```

```
                    'newton-cholesky',
```

```

'sag',
                                     'saga']})

rd.best_params_

{'solver': 'newton-cg', 'penalty': None, 'multi_class': 'ovr'}

from sklearn.linear_model import LogisticRegression
lr1 = LogisticRegression(solver='lbfgs', penalty = None, multi_class
= 'multinomial')
lr1.fit(x_train,y_train)

LogisticRegression(multi_class='multinomial', penalty=None)

lr1.score(x_test,y_test),lr1.score(x_train,y_train)

(0.9187817258883249, 0.9453621346886912)

#From the above result we can say that the model is over-fitted

#Lets check The Cross Validation technique

from sklearn.model_selection import KFold, cross_val_score
kf = KFold(n_splits=5, shuffle=True, random_state=42)
# Cross-validation scores
scores = cross_val_score(lr1, x_train, y_train, cv=kf,
scoring='accuracy')
print("K-Fold Accuracy for training set:", scores.mean())
scores = cross_val_score(lr1, x_test, y_test, cv=kf,
scoring='accuracy')
print("K-Fold Accuracy fpr testing set:", scores.mean())

K-Fold Accuracy for training set: 0.9377328065790534
K-Fold Accuracy fpr testing set: 0.9188461538461539

#Before moving forward lets check the accuracy score for Stratified
Sampling

#Using the Stratified Sampling

df_train , df_test = train_test_split(credit_df,stratify =
credit_df['Class'],test_size = 0.2, random_state = 42)
x_train= df_train.iloc[:, :-1]
y_train = df_train['Class']
x_test= df_test.iloc[:, :-1]
y_test = df_test['Class']
lr2 = LogisticRegression()
lr2.fit(x_train,y_train)

LogisticRegression()

lr2.score(x_test,y_test)

```

```
0.9990519995786665
```

```
lr2.score(x_train,y_train)
```

```
0.9989993197129627
```

```
#So we get a good accuracy by just using the Stratified Sampling
```